

Aleksander Boida

train_acc

Legend:
 - cnn-baseline-1r0.001 (red line)
 - cnn-dropout_0.0_0.3 (dark blue line)
 - cnn-dropout_0.0_0.5 (teal line)
 - cnn-dropout_0.3_0.0 (purple line)
 - cnn-dropout_0.3_0.3 (yellow line)
 - cnn-dropout_0.5_0.5 (light blue line)

val_acc

Legend:
 - cnn-baseline-1r0.001 (red line)
 - cnn-dropout_0.0_0.3 (dark blue line)
 - cnn-dropout_0.0_0.5 (teal line)
 - cnn-dropout_0.3_0.0 (purple line)
 - cnn-dropout_0.3_0.3 (yellow line)
 - cnn-dropout_0.5_0.5 (light blue line)

The figure consists of two line plots. The top plot, titled 'train_acc', shows training accuracy over time. The bottom plot, titled 'val_acc', shows validation accuracy over time. Both plots share the same x-axis, 'Time (minutes)', ranging from 0 to 10. The y-axis for 'train_acc' ranges from 0.7 to 0.9, and for 'val_acc' from 0.76 to 0.82. Four models are compared: 'cnn-baseline-lr0.001' (pink line), 'cnn-maxpool1' (teal line), 'cnn-maxpool2' (light green line), and 'cnn-maxpool3' (brown line). In the 'train_acc' plot, all models show an upward trend, with 'cnn-baseline-lr0.001' reaching the highest accuracy of approximately 0.91 at 6.5 minutes. In the 'val_acc' plot, 'cnn-maxpool2' generally achieves the highest validation accuracy, peaking at about 0.825 around 4 minutes. A specific data point for 'cnn-maxpool1' is highlighted with a box: 11.126; 0.8164.

Time (minutes)	cnn-baseline-lr0.001 (train_acc)	cnn-maxpool1 (train_acc)	cnn-maxpool2 (train_acc)	cnn-maxpool3 (train_acc)	cnn-baseline-lr0.001 (val_acc)	cnn-maxpool1 (val_acc)	cnn-maxpool2 (val_acc)	cnn-maxpool3 (val_acc)
0.5	0.71	0.68	0.70	0.70	0.77	0.75	0.77	0.77
1.0	0.78	0.76	0.78	0.78	0.80	0.79	0.80	0.80
2.0	0.82	0.80	0.81	0.81	0.81	0.80	0.81	0.81
3.0	0.84	0.81	0.82	0.82	0.80	0.81	0.82	0.81
4.0	0.86	0.82	0.84	0.83	0.81	0.81	0.82	0.81
5.0	0.88	0.83	0.85	0.84	0.79	0.81	0.82	0.81
6.0	0.90	0.84	0.86	0.85	0.79	0.81	0.82	0.81
7.0	0.91	0.84	0.86	0.85	0.78	0.81	0.82	0.81
8.0	-	0.84	0.86	0.85	0.79	0.81	0.82	0.81
9.0	-	0.85	0.86	0.85	0.79	0.81	0.82	0.81
10.0	-	0.85	0.86	0.85	0.79	0.81	0.82	0.81

train_acc

— cnn-baseline-lr0.001 — cnn-batchnorm1 — cnn-batchnorm2 — cnn-batchnormfull

Time (seconds)

val_acc

— cnn-baseline-lr0.001 — cnn-batchnorm1 — cnn-batchnorm2 — cnn-batchnormfull

Time (seconds)

The figure consists of two line plots. The top plot, titled 'train_acc', shows training accuracy over time (0 to 21 minutes). The bottom plot, titled 'val_acc', shows validation accuracy over time (0 to 21 minutes). Both plots compare three models: 'cnn-baseline-lr0.001' (red line), 'cnn-maxpool2-dropout_0.3' (blue line), and 'cnn-maxpool2-dropout_0.5' (orange line). In the training accuracy plot, the baseline model shows a steady increase, while the other two models show a rapid initial increase followed by a plateau. In the validation accuracy plot, the baseline model shows a peak around 3 minutes followed by a decline, while the other two models show a steady increase followed by a slight decline.

Time (minutes)	cnn-baseline-lr0.001	cnn-maxpool2-dropout_0.3	cnn-maxpool2-dropout_0.5
0	0.68	0.68	0.68
1	0.78	0.75	0.75
2	0.82	0.78	0.78
3	0.85	0.80	0.80
4	0.87	0.81	0.81
5	0.88	0.82	0.82
6	0.89	0.825	0.825
7	0.90	0.83	0.83
8	-	0.835	0.835
9	-	0.835	0.835
10	-	0.835	0.835
11	-	0.835	0.835
12	-	0.835	0.835
13	-	0.835	0.835
14	-	0.835	0.835
15	-	0.835	0.835
16	-	0.835	0.835
17	-	0.835	0.835
18	-	0.835	0.835
19	-	0.835	0.835
20	-	0.835	0.835
21	-	0.835	0.835

Time (minutes)	cnn-baseline-lr0.001	cnn-maxpool2-dropout_0.3	cnn-maxpool2-dropout_0.5
0	0.77	0.77	0.77
1	0.80	0.80	0.80
2	0.81	0.82	0.81
3	0.81	0.82	0.81
4	0.80	0.82	0.81
5	0.79	0.82	0.81
6	0.79	0.82	0.81
7	0.785	0.825	0.815
8	-	0.825	0.815
9	-	0.825	0.815
10	-	0.825	0.815
11	-	0.825	0.815
12	-	0.825	0.815
13	-	0.825	0.815
14	-	0.825	0.815
15	-	0.825	0.815
16	-	0.825	0.815
17	-	0.825	0.815
18	-	0.825	0.815
19	-	0.825	0.815
20	-	0.825	0.815
21	-	0.825	0.815